Model Accuracy

```
Accuracy with typos added: 78.8%

Threshold | auto-handled | accuracy on those
0.5       | 83%          | 84.1%
0.6       | 65%          | 89.2%
0.7       | 48%          | 93.0%
0.8       | 32%          | 95.4%
0.9       | 14%          | 97.9%

Baseline (always guess the most common): 62.7%
Model accuracy: 78.8%
              precision    recall  f1-score   support

    negative       0.89      0.85      0.87      1835
     neutral       0.59      0.69      0.63       620
    positive       0.73      0.69      0.71       473

    accuracy                           0.79      2928
   macro avg       0.73      0.74      0.74      2928
weighted avg       0.80      0.79      0.79      2928
```

Results

```
'My flight was cancelled and nobody will help me' -> negative (74%)
'Thanks for the quick rebooking, great crew' -> positive (100%)
'Is the gate changed for flight 220?' -> negative (64%)
```

Changed the noise from `rate=0.15` to `0.30`. Re-run Cell 2 and Cell 5, and note how accuracy changes. Changed THRESHOLD to `0.8`, re-run Cell 6, and see which messages now go to a human.

```
'Where to get my boarding pass? ' -> neutral (57%)  -> low confidence: send to a human
'Idk why I chose this airline.' -> negative (60%)  -> low confidence: send to a human
'Can you please tell me where to find the lounge?' -> neutral (70%)  -> low confidence: send to a human
```


```
'Where to get my boarding pass? ' -> neutral (57%)  -> low confidence: send to a human
'Idk why I chose this airline.' -> negative (60%)
'Can you please tell me where to find the lounge?' -> neutral (70%)
```